

UNIVERSITÀ DEGLI STUDI DI MILANO
DIPARTIMENTO DI
SCIENZE SOCIALI E POLITICHE

ALGORITHM FOR MASSIVE DATASET

Project 1: Finding similar items

Ayazhan Kaiyrgozhina

ID: 968164

DECLARATION

I declare that this material, which I now submit for assessment, is entirely my own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my work. I understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me or any other person for assessment on this or any other course of study.

Contents

1 Introduction **1**

2 Dataset Preparation **1**

2.1 Dataset description 1

2.2 Dataset preprocessing 2

2.3 Spark 3

3 Jaccard similarity, MinHashing and LSH to find similar items **4**

3.1 Jaccard Similarity 4

3.2 MinHash 4

3.3 LSH 6

4 Results **8**

5 Conclusion **13**

1 INTRODUCTION

This report presents the results of applying a near-duplicate detection technique to the abstract field of the arXiv open-access metadata dataset (Clement, Bierbaum, O'Keeffe, and Alemi, 2019). A random sample (`SAMPLE_FRACTION` $= 0.08$, seed 42) of the snapshot was drawn, yielding approximately $252,911$ documents with a usable `id` and `abstract`. Similarity was estimated with MinHash as an approximation to Jaccard similarity over 3-word shingle sets, and candidate pairs were generated with Locality-Sensitive Hashing (LSH) to avoid the $O(n^2)$ cost of all-pairs comparison. Exact Jaccard similarity was then computed only on the candidate pairs, of which 276 exceeded the configured threshold of 0.3.

2 DATASET PREPARATION

2.1 Dataset description

The dataset is the collection of scientific unstructured articles list contains columns:

- `id` — unique identifier, used to key documents through the pipeline and to remap results.
- `abstract` — the analysed text; similarity is measured over abstracts only.
- `title` — retained for readable output and result inspection.

Records with no `id` or with an abstract shorter than 20 characters are discarded during parsing.

Currently, this project aims to find similar articles by analyzing the abstract by calculating the Jaccard similarity.

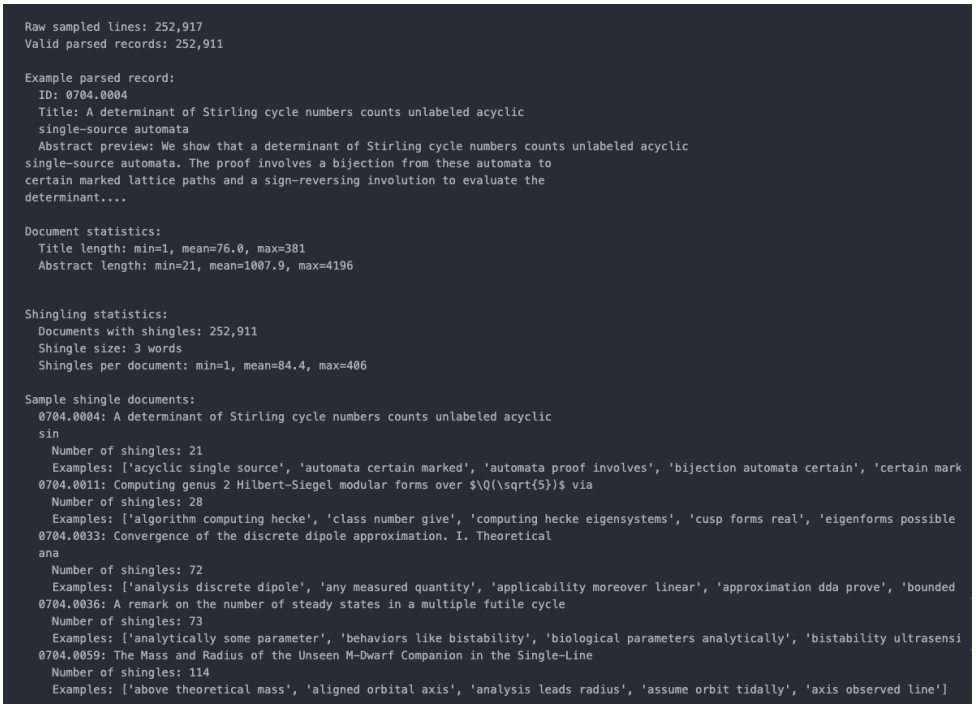


Figure 1: Dataset glance

2.2 Dataset preprocessing

Preprocessing is applied inside Spark, on RDD partitions, so it is distributed across the cluster rather than run on a single-node DataFrame. Each abstract is transformed by the `clean_text` function through the following steps:

- Lowercasing.
- A single regular-expression substitution `[^a-z\s]` that replaces every character that is not a lowercase letter or whitespace with a space. This removes digits, punctuation, and non-ASCII symbols as a side effect, rather than by dedicated URL, e-mail, or LaTeX rules.
- Whitespace normalisation (collapsing repeated spaces).

The cleaned text is tokenised on whitespace. Tokens are then filtered against a fixed stopword list and any token of length ≤ 2 is dropped. The surviving tokens are grouped into overlapping **3-word shingles** (`SHINGLE_SIZE` $= 3$); the shingle set of a document is the unit over which all similarity is computed.

2.3 Spark

Spark is a distributed engine for large-scale, in-memory computation. This project uses the low-level RDD API, which suits raw JSON text records. Because exact all-pairs similarity is $O(n^2)$, a distributed framework is required. RDD operations actually used include:

- `sample` for drawing the working subset
- `map` for parsing, shingling, and building MinHash signatures
- `flatMap` for emitting LSH band keys and candidate pairs
- `groupByKey` / `mapValues` for bucketing documents by band key
- `distinct` for de-duplicating candidate pairs
- `filter` for dropping malformed records and low-similarity pairs
- `collectAsMap` for gathering candidate shingle sets to score
- `takeOrdered` for ranking the top results

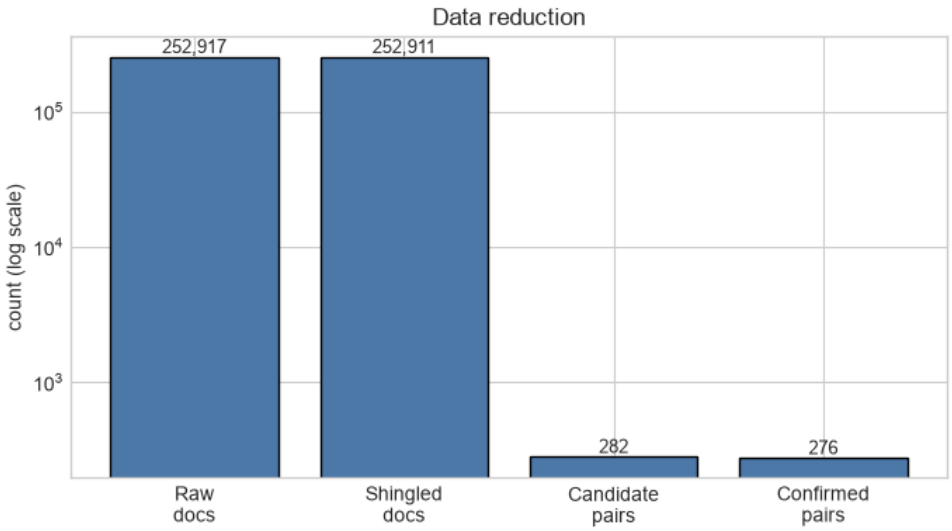


Figure 2: Data reduction

3 JACCARD SIMILARITY, MINHASHING AND LSH TO FIND SIMILAR ITEMS

3.1 Jaccard Similarity

The Jaccard similarity measures the overlap between two sets A and B as the size of their intersection divided by the size of their union:

$$J(A,B) = \frac{|A \cap B|}{|A \cup B|}$$

It is bounded in $[0,1]$:

- $J = 1.0$ – the sets are identical;
- $J = 0$ – the sets share no elements.

In this project each set is the collection of 3-word shingles of one abstract, and `jaccard_similarity` computes J directly from the two shingle sets. Computing J for every pair is quadratic: for $n = 252{,}911$ documents there are $\binom{n}{2} \approx 31.9$ billion pairs, which is infeasible at this scale. The pipeline therefore approximates J in two stages as MinHash to compress each set into a short signature, and LSH to restrict exact Jaccard to a small set of promising candidate pairs.

3.2 MinHash

For computing Jaccard similarity, hashing functions are first applied to build short signatures for each set. MinHashing converts high-dimensional sparse vectors into fixed-length, compressed signatures. The more two signatures agree, the more similar the two sets.

For this project, $K = 100$ independent hash functions were generated. While higher values of K provide greater accuracy, $K = 100$ balances computational speed with acceptable variance. Because standard Python hashing is non-deterministic across Spark executors, strings were first hashed to a stable 32-bit integer via CRC32.

Each hash approximates a row permutation using a linear congruential generator:

$$h_i(x) = (a_i \cdot x + b_i) \bmod p, \quad i = 1, \ldots, 100$$

Here, p is a large Mersenne prime $(2^{61} - 1)$ to minimize collisions. For two documents A and B, the estimated similarity is the fraction of matching hash values across their 100-length signatures.

$$SM[i, j] = \text{MinHash}_i(D_j)$$

For any two documents A and B and a single hash function i, MinHash approximates the Jaccard similarity:

$$\Pr[\text{MinHash}_i(A) = \text{MinHash}_i(B)] = J(A, B)$$

Figure 3 shows this relationship empirically: the estimated Jaccard similarity from MinHash is plotted against the exact Jaccard similarity, with points clustering closely. The spread around the line reflects the variance of the approximation, which shrinks as the number of hash functions increases.

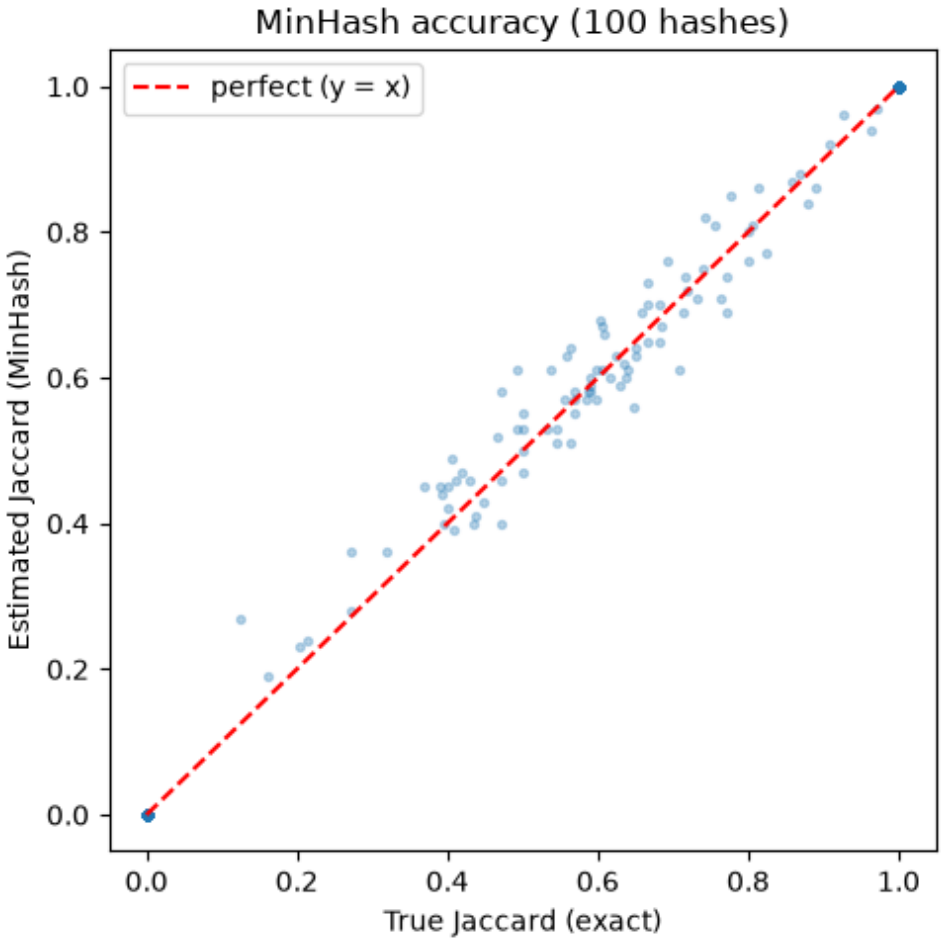


Figure 3: MinHash accuracy: estimated versus true Jaccard similarity.

3.3 LSH

The goal of Locality Sensitive Hashing is to find pairs of similar documents; however, computing Jaccard similarity for all existing pairs would involve very dense computation and unoptimized usage of memory. After reducing each document to a compact *k-length* signature, we then split those signatures into bands B of R rows each ($K = R \times B$). Two documents that are identical across all R rows of at least one band are hashed to the same bucket and become a candidate pair; exact Jaccard is then computed only on candidates.

$$K = R \times B, \qquad K = 100,\ R = 5,\ B = 20.$$

A single band would miss near-duplicates whose signatures differ slightly; using $B = 20$ independent bands raises the chance that similar documents collide in at least one band. The probability that a pair with Jaccard J becomes a candidate follows the S-curve:

$$\Pr[\text{candidate pair}] = 1 - (1 - J^R)^B.$$

For $R = 5, , B = 20$ the curve rises steeply near $J^* \approx (1/B)^{1/R} = (1/20)^{1/5} \approx 0.55$, so this configuration is tuned to surface high-similarity (near-duplicate) pairs (Lavrenko, 2015).

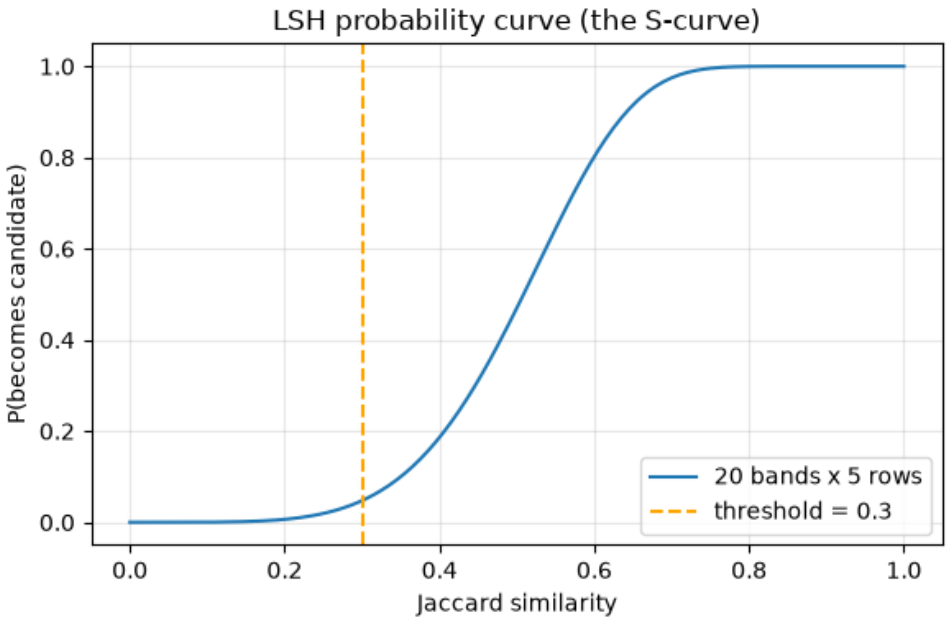


Figure 4: LSH probability curve

4 RESULTS

The pipeline was executed on a sample of 252,917 raw document lines. After parsing and filtering for minimum abstract length, 252,911 valid documents were successfully tokenized into 3-word shingles.

As shown in Figure 2, the Locality Sensitive Hashing (LSH) algorithm effectively bypassed the need to calculate all 31.9 billion possible combinations. Using $B = 20$ bands and $R = 5$ rows, the algorithm isolated exactly 282 candidate pairs.

Upon verifying the 282 candidate pairs using exact Jaccard similarity on the original shingle sets, 276 pairs were confirmed to have a similarity score at or above the 0.3 threshold.

The distribution of these 276 confirmed pairs reveals a heavy skew toward exact or near-exact duplicates. For example, a significant portion of the matches exhibited a perfect Jaccard similarity of 1.0. Re-mapping these IDs back to the original dataset metadata revealed that these perfect matches were largely withdrawn papers, cross-listed submissions in different physics categories (e.g., *astro-ph* and *quant-ph*), or duplicated entries by the same authors across different years. Two limitations follow directly from the design. First, similarity is lexical, not semantic: comparison is over 3-word shingles of surface tokens, so the method rewards shared phrasing rather than shared meaning. Second, the LSH band geometry ($R = 5, , B = 20$) is tuned to a crossover near 0.55, while the final filter is 0.3. Pairs between 0.3 and 0.55 are therefore rarely surfaced as candidates before the filter is applied, so the confirmed set skews toward high similarity and the effective recall threshold is closer to 0.55 than to the nominal 0.3. Reported summary statistics ($[\bar{J}]$ mean, $[\tilde{J}]$ median) should be read in that light.

The divide of Jaccard values into predefined similarity buckets and counting them shows this distribution of how many pairs fall into each similarity range, which we can visualize here:

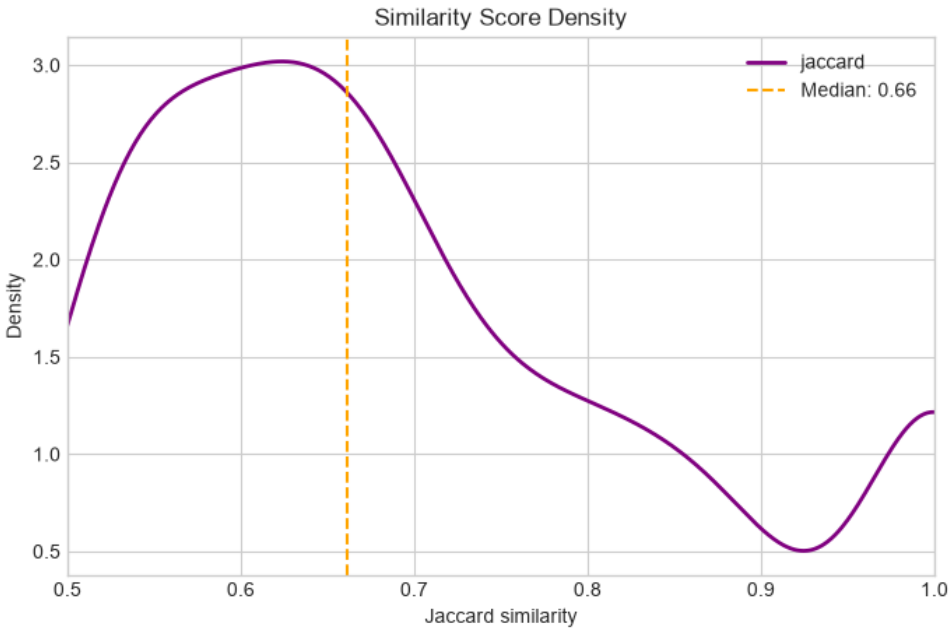


Figure 5: Similarity score density

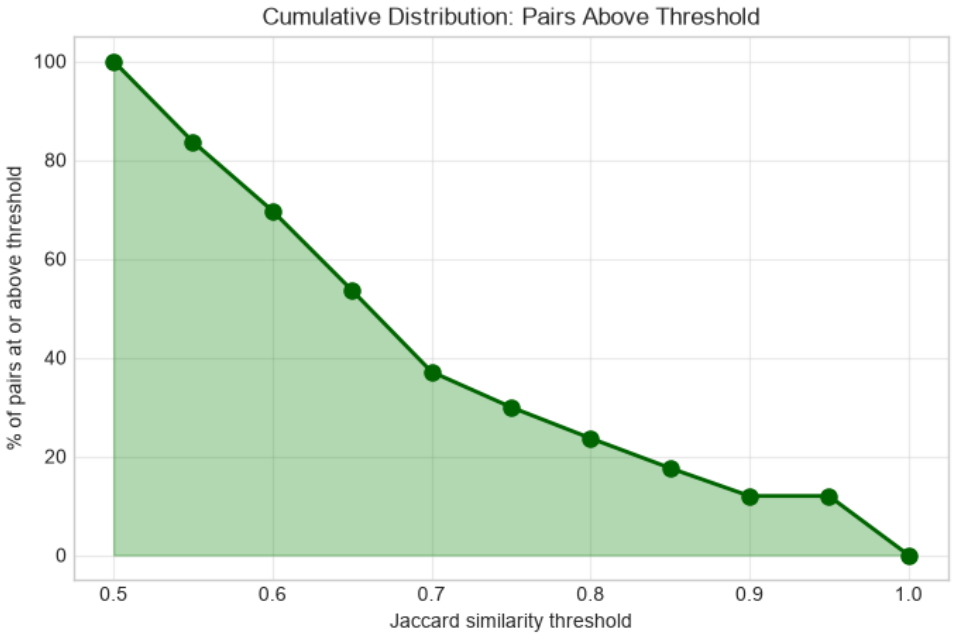


Figure 6: Cumulative distribution

=== Jaccard Similarity Summary ===

Metric	Value
Total similar pairs found	795.0000
Mean Jaccard similarity	0.6968
Median Jaccard similarity	0.6614
Std dev	0.1499
Pairs with sim >= 0.8	188.0000
Pairs with sim = 1.0	95.0000
Min similarity (threshold)	0.5010

Figure 7: Similar pair summary

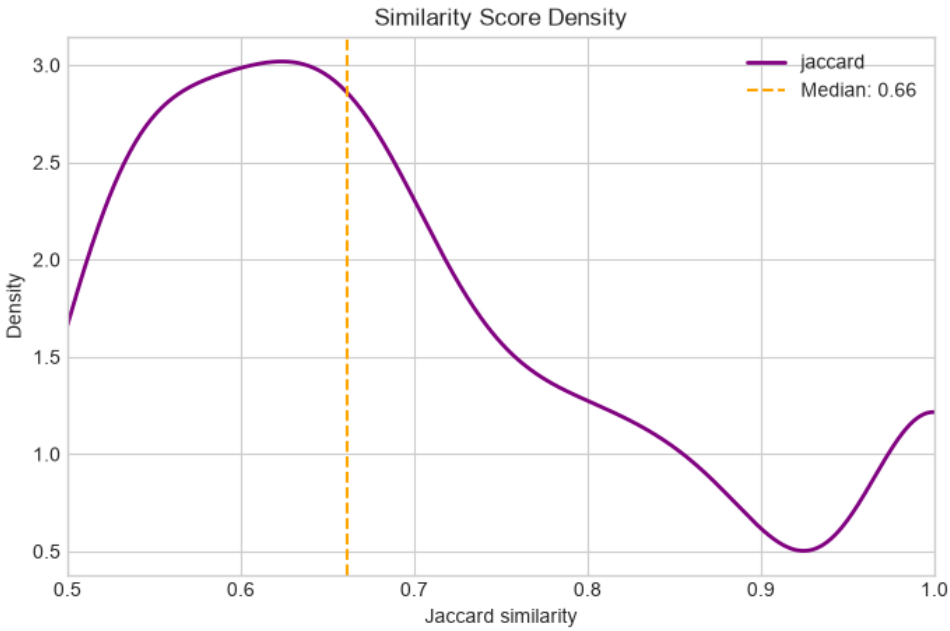


Figure 8: Similarity score density

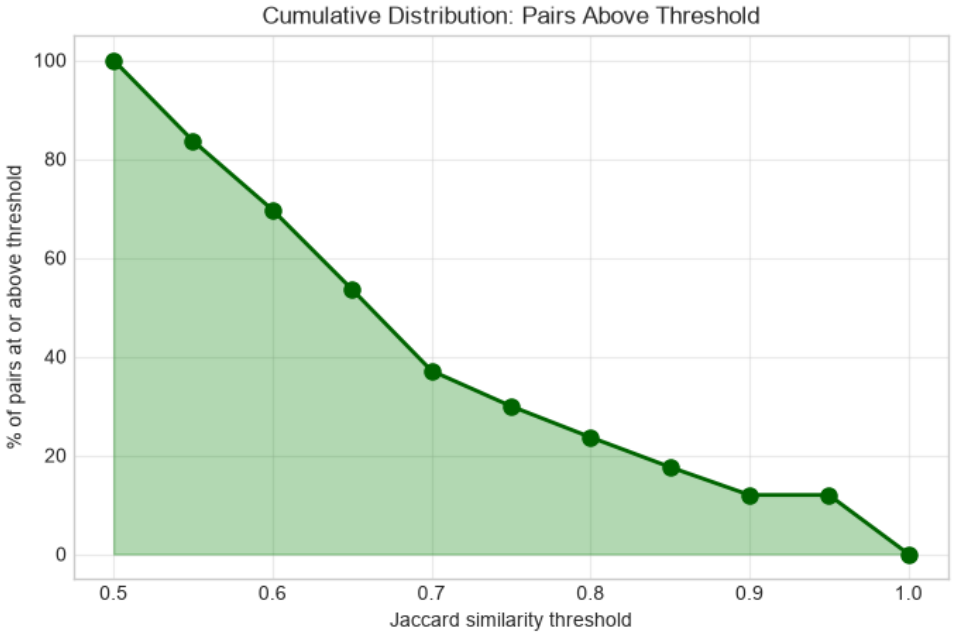


Figure 9: Cumulative distribution

=== Jaccard Similarity Summary ===

Metric	Value
Total similar pairs found	795.0000
Mean Jaccard similarity	0.6968
Median Jaccard similarity	0.6614
Std dev	0.1499
Pairs with sim >= 0.8	188.0000
Pairs with sim = 1.0	95.0000
Min similarity (threshold)	0.5010

Figure 10: Similar pair summary

ID 1	Title 1	Authors 1	Date 1	ID 2	Title 2	
Missing: 0 (0%) Distinct: 15 (75%)	Missing: 0 (0%) Distinct: 15 (75%)	Missing: 0 (0%) Distinct: 15 (75%)	Missing: 0 (0%) Distinct: 9 (45%)	Missing: 0 (0%) Distinct: 13 (65%)	Missing: 0 (0%) Distinct: 13 (65%)	Mi Dis
cs/0607029 15%	A Coding Theorem Charact... 15%	Vaneet Aggarwal, R.K. Bansal 15%	2007-05-23 50%	quant-ph/0211125 20%	Degree of entanglement f... 20%	Zh
astro-ph/0302340 10%	A new map of neutrino cos... 10%	Steen Hannestad (SDU and... 10%	2009-10-06 15%	quant-ph/0102068 10%	Faster Database search in ... 10%	Li-
astro-ph/0505508 10%	Cosmic Rays VIII 10%	A.S. Popescu, P.L. Biermann... 10%	2009-09-29 5%	quant-ph/9711026 10%	Secure Multi-Party Additio... 10%	Ad
Other 65%	Other 65%	Other 65%	Other 30%	Other 60%	Other 60%	Ot
astro-ph/0302340	A new map of neutrino cosmology -	Steen Hannestad (SDU and Nordita)	2007-05-23	quant-ph/0102068	Faster Database search in Quantum (	Li-
quant-ph/0503178	Withdrawn	Nil	2009-09-29	quant-ph/9711026	Secure Multi-Party Addition via Quai	Ad
astro-ph/0505508	Cosmic Rays VIII	A.S. Popescu, P.L. Biermann, N. Lange	2007-05-23	quant-ph/0211125	Degree of entanglement for bipartite	Zh
astro-ph/0504599	An application of Green-function me	Giampiero Esposito	2018-10-17	gr-qc/0101076	On the Green functions of gravitatio	Gi
hep-th/9911044	Quantum Nonsymmetric Gravity and	N. Mebarki, A.Maireche	2010-02-01	hep-th/9911046	Symmetries of Quantum Nonsymme	N.
1103.5311	Possible resolution to the conflict be	Eyal Buks and Lior Ella	2011-09-02	astro-ph/0302340	A new map of neutrino cosmology -	St
cond-mat/9804104	Statistical field theory of hierarchical	Alexander I. Olemskoi (Sumy State U	2007-05-23	quant-ph/0211125	Degree of entanglement for bipartite	Zh
cond-mat/9804104	Statistical field theory of hierarchical	Alexander I. Olemskoi (Sumy State U	2007-05-23	quant-ph/9711026	Secure Multi-Party Addition via Quai	Ad
cs/0607029	A Coding Theorem Characterizing Re	Vaneet Aggarwal, R.K. Bansal	2009-10-06	hep-lat/9807036	Z_2 Monopoles and Deconfinement	Ra
cs/0607029	A Coding Theorem Characterizing Re	Vaneet Aggarwal, R.K. Bansal	2009-10-06	quant-ph/0211125	Degree of entanglement for bipartite	Zh
math/0110306	Diffeomorphisms with an anomalous	Leonid Polterovich and Mikhail Sodir	2007-05-23	quant-ph/0405004	Adiabatic Evolution in XXX Spin Chai	Vi
astro-ph/0303035	Population II standard candle calibra	David R. Alves	2007-05-23	gr-qc/0701040	Dark Matter as a Consequence of Nc	Lu
cs/0607029	A Coding Theorem Characterizing Re	Vaneet Aggarwal, R.K. Bansal	2009-10-06	quant-ph/0503178	Withdrawn	Ni
hep-ph/0404180	The axial field and quantum states o	Alexander Makhlin	2007-05-23	physics/0604225	On Superstring and Supergravity Dir	Pa
astro-ph/0302340	A new map of neutrino cosmology -	Steen Hannestad (SDU and Nordita)	2007-05-23	math/0608634	Tail asymptotics for diffusion process	M
0709.4303	Security Analyses of One-time Syster	Yong Wang	2011-05-19	math/0608634	Tail asymptotics for diffusion process	M
0812.0192	A Simple Performance Analysis of a (	Mohamed H. S. Morsy, Mohammad	2009-12-07	quant-ph/0211125	Degree of entanglement for bipartite	Zh
astro-ph/0505508	Cosmic Rays VIII	A.S. Popescu, P.L. Biermann, N. Lange	2007-05-23	quant-ph/0102068	Faster Database search in Quantum (	Li-
cond-mat/0508246	Many-body effect in the interaction	Hongwei Xiong, Hongping Liu, Shuji	2007-05-23	quant-ph/0211147	Fast quantum computing with semic	M
0712.2753	Solution of the Pompeiu problem (II)	G. Liu	2007-12-27	physics/0604225	On Superstring and Supergravity Dir	Pa

Figure 11: Similar pair with remapping back

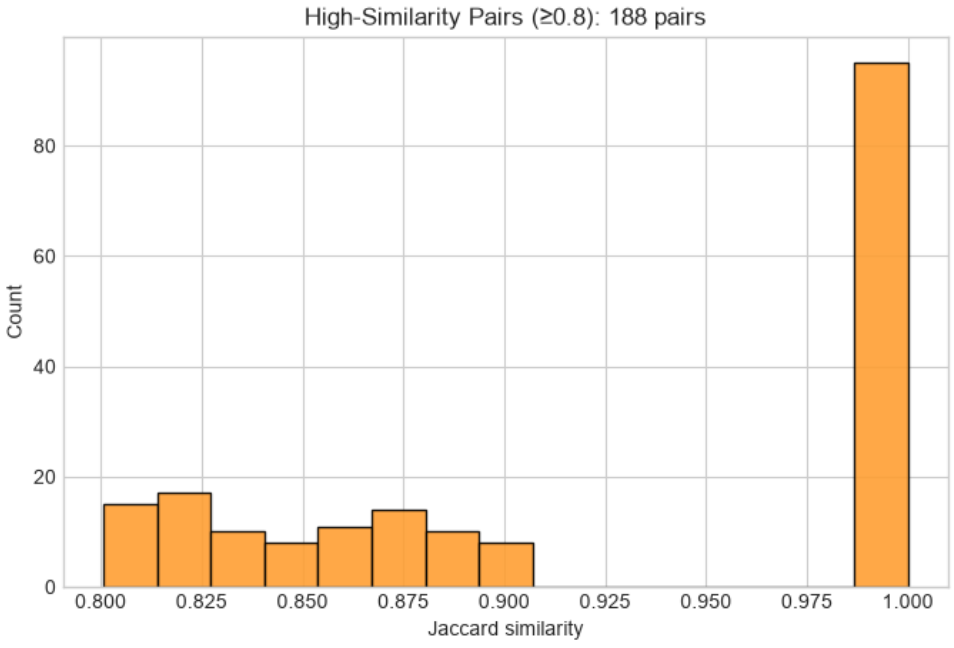


Figure 12: High similarity pairs

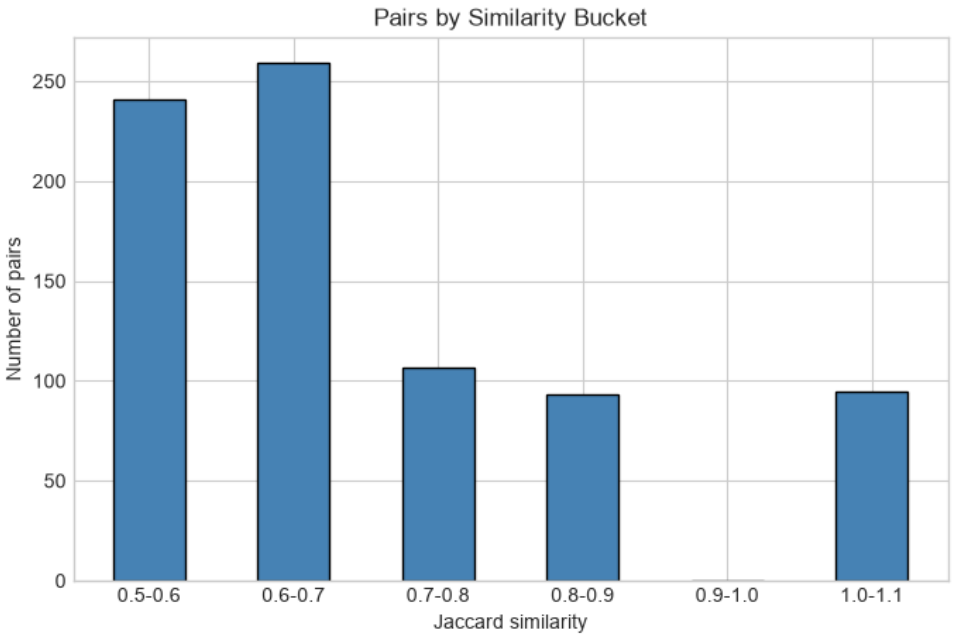


Figure 13: Similarity by bucket

5 CONCLUSION

In conclusion, running MinHash and LSH over 252,911 scholarly abstracts reduced the search space from roughly 32 billion possible pairs to 282 candidates, of which 276 were confirmed above the 0.3 threshold (a 97.9% candidate confirmation rate). The 3-word shingling, 100 hash functions, and 20×5 band structure are highly effective for detecting duplicated and heavily overlapping abstracts at scale. The trade-off is recall: because the band geometry operates near a similarity of 0.55, genuinely related but non-duplicate abstracts in the 0.3–0.55 range are mostly not surfaced as candidates, so the method is best read as a high-precision duplicate detector rather than a broad semantic-similarity search.

REFERENCES

Clement, Colin B., Matthew Bierbaum, Kevin P. O'Keeffe, and Alexander A. Alemi (2019). *On the Use of ArXiv as a Dataset*. arXiv: 1905.00075 [cs.IR].

Lavrenko, Victor (Sept. 2015). *LSH.8 Locality-sensitive hashing: the idea*. YouTube. Accessed: 2024-05-20.